

# STA 6133: Homework 1

Ricardo Cortez & Ben Graf

Due 9 Feb 2021

## 1.

Consider approximating the function  $\sin(x)$  by its Taylor polynomial of degree  $2n - 1$ :

$$p_n(x) = x - \frac{x^3}{3!} + \frac{x^5}{5!} + \cdots + (-1)^{n+1} \frac{x^{2n-1}}{(2n-1)!}$$

### (a)

Write an R function to compute  $p_n(x)$  for  $x \in \mathbb{R}$  and  $n \in \mathbb{N}$ .

```
sinapprox <- function(x, n) {  
  ## sin(x) approximation using Taylor polynomial of degree 2n-1  
  taylor <- 0  
  for (i in 1:n) {  
    taylor <- taylor + (-1)^(i+1) * x^(2*i-1) / factorial(2*i-1)  
  }  
  return(taylor)  
}
```

### (b)

Use your function to approximate  $\sin(x)$  for  $x = \pi$  and  $5\pi$ , and  $n = 10, 20$ , and  $100$ . What lesson(s) can you draw from this example?

```
sin(pi)
```

```
## [1] 1.224647e-16
```

```
sin(5*pi)
```

```
## [1] 6.123234e-16
```

```
sinapprox(pi,10)
```

```
## [1] -5.289183e-10
```

```
sinapprox(pi,20)
```

```
## [1] 3.328057e-16
```

```
sinapprox(pi,100)
```

```
## [1] 3.328057e-16
```

```
sinapprox(5*pi,10)
```

```
## [1] -169520.9
```

```
sinapprox(5*pi,20)
```

```
## [1] -0.2886023
```

```
sinapprox(5*pi,100)
```

```
## [1] -4.578756e-11
```

Because the Taylor polynomial contains numerous terms with negative signs, there is a risk of subtractive cancellation. The true values of  $\sin(\pi)$  and  $\sin(5\pi)$  should be 0. Even the built-in R functions do not provide exactly 0, though their results are on the order of  $10^{-16}$ , so very small. We get similarly small numbers for our Taylor approximations of  $\sin(\pi)$ , though the version with  $n = 10$  is not quite as small as the others. The problems really kick in with approximating  $\sin(5\pi)$ . Because the numbers being subtracted are larger, we see major errors in the approximations. For  $n = 10$ , the approximation is roughly  $-169,000$ , nowhere near 0! With  $n = 20$ , it improves but is still around  $-0.29$ . Finally, with  $n = 100$ , the approximation nears 0. The big lesson here is to beware of Taylor expansions with alternating positive and negative terms!

## 2.

On your computer, find a root for the quadratic equation  $ax^2 + bx + c = 0$  using the formula

$$x = \frac{-b + \sqrt{b^2 - 4ac}}{2a}$$

with the values  $b = 1$  and  $a = c = 10^{-n}$ . Assess the accuracy of your answers for  $n = 1, 5, 10$  by seeing how closely the equation is satisfied. Now think of a clever way to avoid subtractive cancellation and improve your answers.

## Solution:

Finding a root in a non-clever way:

```
# create a quadratic function since we'll be testing several n's and potentially different b's
quad <- function(n,b){
  a <- 10^-n
  c <- a

  root_num_right <- sqrt(b^2-(4*a*c))
  #splitting the equations into separate terms to better understand the subtractive cancellation
  root_num <- -b +root_num_right
  root_den <- 2*a

  quad_root <- root_num/root_den
  true_root <- Re(polyroot(c(c,b,a)))[1] # only take the negative root since that's what we're dealing with

  #use Re to only take the real root of the complex number, checked them all, no imaginary parts.

  df <- data.frame("TrueRoot" =true_root, "QuadRoot" = quad_root)

  #calculate Relative Error
  rel_err <- abs((true_root-quad_root)/true_root)
  df$rel_err <- rel_err
  cat(paste0("Relative Error for a=",a, ", b=",b," ", c=",c, ", and n=",n," is ", rel_err))

  return(df)
}

df1 = quad(1,1)
```

```
## Relative Error for a=0.1, b=1, c=0.1, and n=1 is 2.06063905221897e-15
```


```
df2 = quad(5,1)
```

```
## Relative Error for a=1e-05, b=1, c=1e-05, and n=5 is 8.26403710463034e-08
```

```
df3 = quad(10,1)
```

```
## Relative Error for a=1e-10, b=1, c=1e-10, and n=10 is 1
```

```
plot(c(1,5,10),c(df1$rel_err,df2$rel_err,df3$rel_err), ylab = "Relative Error", xlab = "N")
```



From the plot above, we can see as  $n$  increases the relative error increases to 1, which indicates that subtractive cancellation is present. A closer observation of the individual terms (not shown) indicated that the most aggressive cancellation occurred in between the square root term and the  $-b$  term. A common way of handling this type of cancellation is to use known mathematical identities to eliminate the need for subtraction between those elements. In this problem we used:

$$x^2 - y^2 = (x - y)(x + y)$$

The derivation is as follows:

For simplicity, focus only on the numerator for now, the denominator can be added later.

$$\sqrt{b^2 - 4ac} - b$$

Let  $z = \sqrt{b^2 - 4ac}$  and multiply by a carefully chosen version of 1:

$$\frac{(z - b)(z + b)}{z + b} = \frac{z^2 - b^2}{z + b}$$

Plugging original values back in:

$$\frac{b^2 - 4ac - b^2}{\sqrt{b^2 - 4ac} + b}$$

Simplifying:

$$\frac{-4ac}{\sqrt{b^2 - 4ac} + b}$$

Adding back the original denominator:

$$\frac{-4ac}{2a * (\sqrt{b^2 - 4ac} + b)}$$

Simplify again to obtain the final equation:

$$\frac{-2c}{(\sqrt{b^2 - 4ac} + b)}$$

Modifying the earlier code:

```
quad_new <- function(n,b){
  a <- 10^-n
  c <- a

  quad_root = -2*c / (sqrt(b^2 - 4*a*c) + b)
  true_root = Re(polyroot(c(c,b,a)))[1] # only take the negative root since that's what we're dealing with

  #use Re to only take the real root of the complex number, checked them all, no imag parts.

  df <- data.frame("TrueRoot" =true_root, "QuadRoot" = quad_root)

  #calculate Relative Error
  rel_err <- abs((true_root-quad_root)/true_root)
  df$rel_err <- rel_err
  cat(paste0("Relative Error for a=",a, ", b=",b,", c=",c, ", and n=",n," is ", rel_err))

  return(df)
}

df1_new <- quad_new(1,1)
```

```
## Relative Error for a=0.1, b=1, c=0.1, and n=1 is 0
```


```
df2_new <-quad_new(5,1)
```

```
## Relative Error for a=1e-05, b=1, c=1e-05, and n=5 is 1.69406589433919e-16
```

```
df3_new <-quad_new(10,1)
```

```
## Relative Error for a=1e-10, b=1, c=1e-10, and n=10 is 1.29246970711411e-16
```

```
plot(c(1,5,10),c(df1_new$rel_err,df2_new$rel_err,df3_new$rel_err), ylab = "Relative Error", xlab = "N",ylim = c(0,1))
```



Using the new solution provides a better estimate of the roots when subtractive cancellation is present in the original equation, it is clear to see in the plot above that the relative error stays near 0 for all  $N$ 's.

### 3.

The following data are an i.i.d. sample from the  $Cauchy(\theta, 1)$  distribution ( $\theta \in \mathbb{R}$ ):

1.77	-0.23	2.76	3.80	3.47	56.75
-1.34	4.24	-2.44	3.29	3.71	-2.40
4.53	-0.07	-1.05	-13.87	-2.53	-1.75
0.27	43.21				

#### (a)

Plot the log-likelihood function. Then find the MLE of  $\theta$  using Newton's method. Try all the following starting values: -11, -1, 0, 1.5, 4, 4.7, 7, 38. Discuss your results. Is the sample mean a good starting point?

The density for  $Cauchy(\theta, \sigma)$  is:

$$f(x | \theta, \sigma) = \frac{1}{\pi\sigma} \frac{1}{1 + \left(\frac{x-\theta}{\sigma}\right)^2}$$

Because  $\sigma = 1$ , we get:

$$f(x | \theta, \sigma) = \frac{1}{\pi} \frac{1}{1 + (x - \theta)^2}, \quad -\infty < \theta < \infty, \quad -\infty < x < \infty$$

The joint density is then:

$$f(\mathbf{x} | \theta, 1) = \prod_{i=1}^n \frac{1}{\pi} \frac{1}{1 + (x_i - \theta)^2} = \pi^{-n} \prod_{i=1}^n \frac{1}{1 + (x_i - \theta)^2}$$

The log-likelihood is then:

$$l(\theta) = -n \ln \pi - \sum_{i=1}^n \ln(1 + (x_i - \theta)^2)$$

Its first derivative is:

$$l'(\theta) = \sum_{i=1}^n \frac{2(x_i - \theta)}{1 + (x_i - \theta)^2}$$

And its second derivative is:

$$l''(\theta) = \sum_{i=1}^n \frac{(1 + (x_i - \theta)^2)(-2) - 2(x_i - \theta)(-2(x_i - \theta))}{(1 + (x_i - \theta)^2)^2}$$



$$= \sum_{i=1}^n \frac{-2(1 + (x_i - \theta)^2) + 4(x_i - \theta)^2}{(1 + (x_i - \theta)^2)^2} = \sum \frac{2(-1 + (x_i - \theta)^2)}{(1 + (x_i - \theta)^2)^2}$$

The corresponding R functions are:

```
#Import the data
data <- c(1.77, -0.23, 2.76, 3.80, 3.47, 56.75, -1.34, 4.24, -2.44, 3.29,
          3.71, -2.40, 4.53, -0.07, -1.05, -13.87, -2.53, -1.75, 0.27, 43.21)

n <- length(data)

#log likelihood cauchy function
log_cauchy <- function(theta){
  return(-n*log(pi)-sum(log(1+(data-theta)^2)))
}

#first derivative of the log-likelihood
log_cauchy_prime <- function(theta){
return(sum(2*(data-theta)/(1+(data-theta)^2)))
}

#second derivative of the log-likelihood
log_cauchy_2prime <- function(theta){
  return(sum(2*(-1+(data-theta)^2) / (1+(data-theta)^2)^2))
}

#function for log-cauchy and log-cauchy prime
cauchy <- function(theta){
  return(c(log_cauchy(theta),
          log_cauchy_prime(theta)))
}


#function for log-cauchy prime and 2nd prime of log-cauchy
cauchy_primes <- function(theta){
  return(c(log_cauchy_prime(theta),log_cauchy_2prime(theta)))
}

theta_list <- seq(-50,50,.1)
values <- sapply(theta_list, cauchy)

log_cauchy_values <- values[1,]
log_cauchy_prime_values <- values[2,]


plot(theta_list,log_cauchy_values, type = "l", main = "Log Cauchy Distribution",xlab
= "Theta", ylab = "log-likelihood")
```

## Log Cauchy Distribution



```
plot(theta_list, log_cauchy_prime_values, type = "l", main = "Log Prime Cauchy Distribution", xlab = "Theta", ylab = "prime-log-likelihood")
abline(h = 0, col = "red")
```

## Log Prime Cauchy Distribution



```

# Newton's Method function provided by Dr. DeOliveira
newtonraphson <- function(ftn, x0, tol = 1e-9, max.iter = 100) {
  ## Newton_Raphson algorithm for solving ftn(x)[1] == 0
  ## It is assumed that ftn is a function of a single variable that returns
  ## the function value and its first derivative as a vector of length 2.
  ## x0 is the initial guess of the root.
  ## The algorithm terminates when the function value is within distance tol of 0,
  ## or the number of iterations exceeds max.iter, whichever happens first
  # initialise
  x <- x0
  fx <- ftn(x)
  iter <- 0
  # continue iterating until stopping conditions are met
  while ((abs(fx[1]) > tol) & (iter < max.iter)) {
    x <- x - fx[1]/fx[2]
    fx <- ftn(x)
    iter <- iter+1
    #cat("At iteration", iter, "value of x is:", x, "\n")
  }
  # output depends on success of algorithm
  if (abs(fx[1]) > tol) {
    #cat("Algorithm failed to converge\n")
    return(NULL)
  } else {
    #cat("Algorithm converged\n")
    return(c(root = x, f.val = fx[1]))
  }
}

```

We test using the recommended starting values:

```

mle_values <- roots <- test_values <- c(-11,-1,0,1.5,4,4.7,7,38, mean(data))

for(i in 0:length(test_values)){
  root <- newtonraphson(cauchy_primes,test_values[i])[1]
  roots[i] <- root
  mle_values[i] <- log_cauchy(root)
}
roots

```

```

## [1] -5.381711e+10 -1.922866e-01 -1.922866e-01  1.713587e+00  2.817472e+00
## [6] -1.922866e-01  4.104085e+01  4.279538e+01  5.487662e+01

```

```

#### max value is at -1
cat(paste0("The maximum log-likelihood value of ", test_values[which.max(mle_values)]
, " occurs when theta = ", roots[which.max(mle_values)]))

```

```
## The maximum log-likelihood value of -1 occurs when theta = -0.192286613229646
```

$$\hat{\theta}_{MLE} = -0.1923.$$

Utilizing different starting points resulted in getting several different root values, which, given the execution pattern of the Newton-Raphson algorithm and the log-prime Cauchy distribution plot above, would be expected. Using the sample mean for a Cauchy distribution is not advantageous because the Cauchy distribution does not have a population mean. In fact, the sample mean as a starting point does *not* lead to the MLE.

(b)

Apply the bisection method with bracketing interval  $[-1, 1]$ . Comment on the result. Now run the bisection method with another bracketing interval that makes the algorithm converge to the ‘wrong root’, i.e., to a local maximum.

```
# Bisection Method function provided by Dr. DeOliveira
bisection <- function(f, x1, x2, maxit = 1000, tol = 1e-7, stop.fval = TRUE) {
  # it uses absolute error as stopping criterion when stop.fval = FALSE
  f1 <- f(x1)
  if (abs(f1) < tol)
    return(x1)
  f2 <- f(x2)
  if (abs(f2) < tol)
    return(x2)
  if (f1 * f2 > 0)
    stop("f has equal sign at endpoints of initial interval")
  if (f1 > 0) { # swap x1 and x2
    tmp <- x1 ; x1 <- x2 ; x2 <- tmp
  }
  n <- 0 # counter
  x <- x1
  repeat {
    n <- n + 1
    x.mid <- (x1 + x2) / 2
    f.mid <- f(x.mid)
    if(stop.fval == TRUE) {
      if(abs(f.mid) < tol | n == maxit)
        break
    }
    else if(abs(x.mid - x) < tol | n == maxit)
      break
    if(f.mid < 0){ x1 <- x.mid ; x <- x1}
    else { x2 <- x.mid ; x <- x2 }
  }
  return(list(root = x.mid, f = f.mid, iter = n))
}

bisection(log_cauchy_prime, -1, 1, tol = 1e-9)
```

```
## $root
## [1] -0.1922866
##
## $f
## [1] 5.166203e-10
##
## $iter
## [1] 31
```

The Bisection took 31 iterations to converge with the given interval. The Newton-Raphson at worst took about the same number of iterations but often took fewer than 10 iterations. Below we see a poor choice of bracketing interval results in converging to a local maximum.

```
bisection(log_cauchy_prime, 1, 2, tol = 1e-9)  #This bracket converges to a local maximum, not the global
```

```
## $root
## [1] 1.713587
##
## $f
## [1] 3.384803e-10
##
## $iter
## [1] 29
```

**(c)**

Apply a fixed-point iteration with starting value  $-1$  and scaling choices  $\alpha = 1, 0.64, 0.25$ . Discuss your results.

```

fixedpoint <- function(f, starting, alpha, maxit = 1000, tol = 1e-9) {
  #  $g(x) = x + \alpha * f(x)$ 
  x <- starting
  fx <- f(starting)
  iter <- 0
  while ((abs(fx) > tol) & (iter < maxit)) {
    x <- x + alpha*fx
    fx <- f(x)
    iter <- iter + 1
    #cat("At iteration", iter, "value of x is:", x, "\n")
  }
  if (abs(fx) > tol) {
    cat("Algorithm failed to converge\n")
    return(NULL)
  } else {
    cat("Algorithm converged\n")
    return(list(root = x, f.val = fx, iter = iter))
  }
}

fixedpoint(log_cauchy_prime, -1, 1)

```

```
## Algorithm failed to converge
```

```
## NULL
```

```
fixedpoint(log_cauchy_prime, -1, 0.64)
```

```
## Algorithm converged
```

```

## $root
## [1] -0.1922866
##
## $f.val
## [1] 9.973122e-10
##
## $iter
## [1] 700

```

```
fixedpoint(log_cauchy_prime, -1, 0.25)
```

```
## Algorithm converged
```

```
## $root
## [1] -0.1922866
##
## $f.val
## [1] 4.31729e-10
##
## $iter
## [1] 16
```

```
fixedpoint(log_cauchy_prime, -1, 0.2)
```

```
## Algorithm converged
```

```
## $root
## [1] -0.1922866
##
## $f.val
## [1] 5.361135e-10
##
## $iter
## [1] 24
```

For an  $\alpha$  of 1, the fixed-point algorithm failed to converge. Reducing  $\alpha$  to 0.64 does achieve convergence, but it takes 700 iterations. Further reducing  $\alpha$  to 0.25 takes a mere 16 iterations! Reducing  $\alpha$  any lower with the same starting point begins increasing the number of iterations again.



## 4.

There were 46 crude oil spills of at least 1000 barrels from tankers in U.S. waters during 1974–1999. The web page <http://www.stat.colostate.edu/computationalstatistics/>

(<http://www.stat.colostate.edu/computationalstatistics/>) contains the following data:

$N_i$  = number of spills in year  $i$ ;

$b_{i1}$  = amount of oil shipped through U.S. waters from import/export shipments in year  $i$  (in billions of barrels [Bbbl]);

$b_{i2}$  = amount of oil shipped through U.S. waters U.S. from domestic shipments in year  $i$ .

Suppose we use the Poisson process model that assumes  $N_i | b_{i1}, b_{i2} \sim \text{Poisson}(\alpha_1 b_{i1} + \alpha_2 b_{i2})$ . The volume of oil shipped is a measure of spill risk, so  $\alpha_1$  and  $\alpha_2$  represent the rates of spill occurrences per Bbbl during import/export and domestic shipments, respectively.

### (a)

Derive Newton's iterative algorithm for finding the MLE of  $\alpha = (\alpha_1, \alpha_2)$ , and write an R function to implement this algorithm. Then run this iterative algorithm to compute the MLE of  $\alpha$  based on the above dataset.

The joint density of  $\mathbf{N}$  is:

$$f(\mathbf{N} | \alpha) = \prod_{i=1}^n e^{-(\alpha_1 b_{i1} + \alpha_2 b_{i2})} \frac{(\alpha_1 b_{i1} + \alpha_2 b_{i2})^{N_i}}{N_i!}$$

The log-likelihood of  $\alpha$  is therefore:

$$l(\alpha) = \sum_{i=1}^n \left( -(\alpha_1 b_{i1} + \alpha_2 b_{i2}) + N_i \ln(\alpha_1 b_{i1} + \alpha_2 b_{i2}) - \ln N_i! \right)$$

Its first derivatives are:

$$\frac{\partial l}{\partial \alpha_1} = \sum_{i=1}^n \left( -b_{i1} + N_i \frac{b_{i1}}{\alpha_1 b_{i1} + \alpha_2 b_{i2}} \right)$$

$$\frac{\partial l}{\partial \alpha_2} = \sum_{i=1}^n \left( -b_{i2} + N_i \frac{b_{i2}}{\alpha_1 b_{i1} + \alpha_2 b_{i2}} \right)$$

The gradient is defined as:

$$\nabla l(\alpha) = \begin{pmatrix} \frac{\partial l}{\partial \alpha_1} \\ \frac{\partial l}{\partial \alpha_2} \end{pmatrix}$$

The second derivatives are:

$$\frac{\partial^2 l}{\partial \alpha_1^2} = \sum_{i=1}^n \frac{-N_i b_{i1}^2}{(\alpha_1 b_{i1} + \alpha_2 b_{i2})^2}$$

$$\frac{\partial^2 l}{\partial \alpha_1 \partial \alpha_2} = \frac{\partial^2 l}{\partial \alpha_2 \partial \alpha_1} = \sum_{i=1}^n \frac{-N_i b_{i1} b_{i2}}{(\alpha_1 b_{i1} + \alpha_2 b_{i2})^2}$$

$$\frac{\partial^2 l}{\partial \alpha_2^2} = \sum_{i=1}^n \frac{-N_i b_{i2}^2}{(\alpha_1 b_{i1} + \alpha_2 b_{i2})^2}$$

The Hessian matrix is defined as:

$$H_l(\boldsymbol{\alpha}) = \begin{bmatrix} \frac{\partial^2 l}{\partial \alpha_1^2} & \frac{\partial^2 l}{\partial \alpha_1 \partial \alpha_2} \\ \frac{\partial^2 l}{\partial \alpha_2 \alpha_1} & \frac{\partial^2 l}{\partial \alpha_2^2} \end{bmatrix}$$

Newton's iterative algorithm is therefore defined as:

$$\boldsymbol{\alpha}_{n+1} = \boldsymbol{\alpha}_n - (H_l(\boldsymbol{\alpha}_n))^{-1} \nabla l(\boldsymbol{\alpha}_n), n = 0, 1, 2, \dots$$

The R functions to implement this are as follows:

```
setwd("/Users/Ben/Library/Mobile Documents/com~apple~CloudDocs/Documents/UTSA Master's/Semester 4/STA 6133 Simulation & Statistical Computing/Homework/Homework 1/")
#readLines("oilspills.dat")
oil <- read.table("oilspills.dat", header = TRUE)

# Log-likelihood function for Poisson
llpois <- function(alphas, bs, N) {
  ## alphas should be 1x2, bs should be 2xn, N should be 1xn
  term <- alphas %*% bs
  return(sum(-term + N*log(term) - log(factorial(N))))
}

# Gradient of log-likelihood for Poisson
gradpois <- function(alphas, bs, N) {
  ## alphas should be 1x2, bs should be 2xn, N should be 1xn
  term <- alphas %*% bs
  drv1 <- sum(-bs[1,] + N*bs[1,]/term)
  drv2 <- sum(-bs[2,] + N*bs[2,]/term)
  return(as.matrix(c(drv1,drv2)))
}

# Hessian of log-likelihood for Poisson
hesspois <- function(alphas, bs, N) {
  ## alphas should be 1x2, bs should be 2xn, N should be 1xn
  term <- alphas %*% bs
  drv11 <- sum(-N * (bs[1,])^2 / (term^2))
  drv12 <- sum(-N * bs[1,] * bs[2,] / (term^2))
}
```

```

drv22 <- sum(-N * (bs[2,])^2 / (term^2))
return(matrix(c(drv11,drv12, drv12,drv22), nrow = 2, ncol = 2, byrow = TRUE))
}

# Newton's Method adapted for this problem
newtonmulti <- function(grad, hess, x0, data1, data2, tol = 1e-9, max.iter = 100) {
  ## Newton_Raphson algorithm for solving grad == 0.
  ## x0 is the initial guess of the roots.
  ## The algorithm terminates when the function value is within distance tol of 0,
  ## or the number of iterations exceeds max.iter, whichever happens first.

  # initialise
  x <- x0
  fx <- grad(x, data1, data2)
  iter <- 0
  # continue iterating until stopping conditions are met
  while ((abs(fx[1]) > tol) | (abs(fx[2]) > tol)) & (iter < max.iter)) {
    x <- x - t(solve(hess(x,data1,data2)) %*% grad(x,data1,data2))
    fx <- grad(x, data1, data2)
    iter <- iter+1
    cat("At iteration", iter, "value of x is:", x, "and f is:", fx, "\n")
  }

  # output depends on success of algorithm
  if ((abs(fx[1]) > tol) | (abs(fx[2]) > tol)) {
    cat("Algorithm failed to converge\n")
    return(NULL)
  } else {
    cat("Algorithm converged\n")
    return(c(root = x, f.val = fx))
  }
}

```

Testing the function with a variety of starting points yields:

```

# Run Newton's Method
testalphas <- t(as.matrix(c(1,1)))
bs <- t(as.matrix(oil[,3:4]))
N <- oil[,2]
res <- newtonmulti(gradpois, hesspois, testalphas, bs, N)

```

```

## At iteration 1 value of x is: 1.090831 0.9426757 and f is: 0.05893873 0.01547945
## At iteration 2 value of x is: 1.097128 0.9375748 and f is: 0.0002198173 5.441683e-
05
## At iteration 3 value of x is: 1.097153 0.9375546 and f is: 3.178962e-09 7.871014e-
10
## At iteration 4 value of x is: 1.097153 0.9375546 and f is: -1.110223e-16 4.440892e-
16
## Algorithm converged

```

```
llpois(res[1:2], bs, N)
```

```
## [1] -48.02716
```

```
newtonmulti(gradpois, hesspois, t(as.matrix(c(0.5,0.5))), bs, N)
```

```
## At iteration 1 value of x is: 0.7727078 0.7356689 and f is: 9.934929 6.243195
## At iteration 2 value of x is: 1.001826 0.8965661 and f is: 2.117605 1.293061
## At iteration 3 value of x is: 1.08905 0.9364765 and f is: 0.1456226 0.08299321
## At iteration 4 value of x is: 1.097096 0.9375682 and f is: 0.0008103152 0.00039828
23
## At iteration 5 value of x is: 1.097153 0.9375546 and f is: 2.760745e-08 1.037014e-
08
## At iteration 6 value of x is: 1.097153 0.9375546 and f is: -1.110223e-16 4.440892e
-16
## Algorithm converged
```

```
##          root1          root2          f.val1          f.val2
## 1.097153e+00  9.375546e-01 -1.110223e-16  4.440892e-16
```

```
newtonmulti(gradpois, hesspois, t(as.matrix(c(0.1,0.5))), bs, N)
```

```
## At iteration 1 value of x is: 0.2055593 0.8153307 and f is: 45.0863 20.6025
## At iteration 2 value of x is: 0.4085579 1.114891 and f is: 17.60121 7.294911
## At iteration 3 value of x is: 0.7192223 1.176014 and f is: 5.309933 1.781389
## At iteration 4 value of x is: 0.9984792 1.019079 and f is: 0.9585817 0.24154
## At iteration 5 value of x is: 1.091235 0.9425422 and f is: 0.05339925 0.012897
## At iteration 6 value of x is: 1.097132 0.9375719 and f is: 0.0001887155 4.666506e-
05
## At iteration 7 value of x is: 1.097153 0.9375546 and f is: 2.344399e-09 5.804666e-
10
## At iteration 8 value of x is: 1.097153 0.9375546 and f is: 9.436896e-16 8.881784e-
16
## Algorithm converged
```

```
##          root1          root2          f.val1          f.val2
## 1.097153e+00  9.375546e-01  9.436896e-16  8.881784e-16
```

```
newtonmulti(gradpois, hesspois, t(as.matrix(c(0.8,0.5))), bs, N)
```

```
## At iteration 1 value of x is: 1.020556 0.7477629 and f is: 3.400149 2.595636
## At iteration 2 value of x is: 1.093548 0.9057989 and f is: 0.3543095 0.2991098
## At iteration 3 value of x is: 1.097232 0.9368358 and f is: 0.005220441 0.005188247
## At iteration 4 value of x is: 1.097153 0.9375543 and f is: 1.525131e-06 1.74987e-06
## At iteration 5 value of x is: 1.097153 0.9375546 and f is: 1.797451e-13 2.146616e-13
## Algorithm converged
```

```
##          root1          root2          f.val1          f.val2
## 1.097153e+00 9.375546e-01 1.797451e-13 2.146616e-13
```

```
newtonmulti(gradpois, hesspois, t(as.matrix(c(2,0.5))), bs, N)
```

```
## At iteration 1 value of x is: 0.5397918 1.334004 and f is: 8.537706 2.489554
## At iteration 2 value of x is: 0.8855011 1.121394 and f is: 2.126835 0.4806032
## At iteration 3 value of x is: 1.070158 0.9606329 and f is: 0.243707 0.05690754
## At iteration 4 value of x is: 1.096722 0.9379148 and f is: 0.003902619 0.0009604976
## At iteration 5 value of x is: 1.097152 0.9375547 and f is: 1.004578e-06 2.487083e-07
## At iteration 6 value of x is: 1.097153 0.9375546 and f is: 6.544765e-14 1.626477e-14
## Algorithm converged
```

```
##          root1          root2          f.val1          f.val2
## 1.097153e+00 9.375546e-01 6.544765e-14 1.626477e-14
```

```
newtonmulti(gradpois, hesspois, t(as.matrix(c(2,2))), bs, N)
```

```
## At iteration 1 value of x is: 0.3633246 -0.2292973 and f is: 19.33009 -145.1933
## At iteration 2 value of x is: 0.6813456 -0.430406 and f is: 3.548708 -76.50709
## At iteration 3 value of x is: 1.20365 -0.7617812 and f is: -3.752098 -41.64224
## At iteration 4 value of x is: 1.912676 -1.215075 and f is: -6.239723 -23.16762
## At iteration 5 value of x is: 2.584328 -1.653179 and f is: -5.39579 -12.03947
## At iteration 6 value of x is: 2.924771 -1.888807 and f is: -2.417535 -4.272862
## At iteration 7 value of x is: 2.998046 -1.947498 and f is: -0.3045471 -0.5066411
## At iteration 8 value of x is: 3.004649 -1.953398 and f is: -0.003872094 -0.006456701
## At iteration 9 value of x is: 3.004737 -1.953477 and f is: -6.418283e-07 -1.071488e-06
## At iteration 10 value of x is: 3.004737 -1.953477 and f is: 4.463097e-14 6.028511e-14
## Algorithm converged
```

##	root1	root2	f.val1	f.val2
##	3.004737e+00	-1.953477e+00	4.463097e-14	6.028511e-14

All but one of these starting points converge to the same root,  $\hat{\alpha}_{MLE} = (1.0972, 0.9376)$ . (The final starting point sees  $\alpha_2$  converge to a negative number, which seems inappropriate given the Poisson distribution.) Plotting the log-likelihood for a range of  $\alpha$ s (not shown) to check that we are converging to a maximum reveals that the maximum does appear to be where both  $\alpha$ s are between 0 and 2, with the highest-valued sample point being (1.1, 0.9), so our root appears to be correct!

## (b)

Derive the Fisher scoring iterative algorithm for finding the MLE of  $\alpha$ , and write an R function to implement this algorithm. Then run this iterative algorithm to compute the MLE of  $\alpha$  based on the above dataset.

The terms of the Fisher Information matrix are:

$$(I(\alpha))_{ij} = -E_{\alpha} \left[ \frac{\partial^2}{\partial \alpha_i \partial \alpha_j} l(\alpha, \mathbf{N}) \right]$$

The first term can therefore be calculated to be:

$$(I(\alpha))_{11} = E \left[ \frac{\partial^2 l}{\partial \alpha_1^2} \right] = E \left[ \sum_{i=1}^n \frac{N_i b_{i1}^2}{(\alpha_1 b_{i1} + \alpha_2 b_{i2})^2} \right] = \sum_{i=1}^n \frac{b_{i1}^2 E[N_i]}{(\alpha_1 b_{i1} + \alpha_2 b_{i2})^2}$$

Because  $E[N_i]$  is the Poisson parameter, we get:

$$= \sum_{i=1}^n \frac{b_{i1}^2 (\alpha_1 b_{i1} + \alpha_2 b_{i2})}{(\alpha_1 b_{i1} + \alpha_2 b_{i2})^2} = \sum_{i=1}^n \frac{b_{i1}^2}{\alpha_1 b_{i1} + \alpha_2 b_{i2}}$$

Similarly:

$$(I(\alpha))_{12} = (I(\alpha))_{21} = \sum_{i=1}^n \frac{b_{i1} b_{i2}}{\alpha_1 b_{i1} + \alpha_2 b_{i2}}$$

$$(I(\alpha))_{22} = \sum_{i=1}^n \frac{b_{i2}^2}{\alpha_1 b_{i1} + \alpha_2 b_{i2}}$$

Running Newton's algorithm for the same set of starting points yields:

```
# Fisher Information of log-likelihood for Poisson
ipois <- function(alphas, bs, N) {
  ## alphas should be 1x2, bs should be 2xn, N should be 1xn
  term <- alphas %*% bs
  drv11 <- -sum((bs[1,])^2 / term)
  drv12 <- -sum(bs[1,] * bs[2,] / term)
  drv22 <- -sum((bs[2,])^2 / term)
  return(matrix(c(drv11,drv12, drv12,drv22), nrow = 2, ncol = 2, byrow = TRUE))
}

# Run Newton's Method
res2 <- newtonmulti(gradpois, ipois, testalphas, bs, N)
```

```
## At iteration 1 value of x is: 1.117785 0.9062845 and f is: -0.05712725 0.07045913
## At iteration 2 value of x is: 1.090702 0.9473314 and f is: 0.01871967 -0.02155272
## At iteration 3 value of x is: 1.099195 0.9344591 and f is: -0.005841277 0.00687103
7
## At iteration 4 value of x is: 1.096508 0.9385308 and f is: 0.001850667 -0.00216218
## At iteration 5 value of x is: 1.097356 0.9372463 and f is: -0.0005835163 0.0006831
982
## At iteration 6 value of x is: 1.097088 0.9376519 and f is: 0.000184263 -0.00021559
48
## At iteration 7 value of x is: 1.097173 0.9375239 and f is: -5.815872e-05 6.806245e
-05
## At iteration 8 value of x is: 1.097146 0.9375643 and f is: 1.835935e-05 -2.148428e
-05
## At iteration 9 value of x is: 1.097155 0.9375515 and f is: -5.795342e-06 6.781905e
-06
## At iteration 10 value of x is: 1.097152 0.9375556 and f is: 1.829394e-06 -2.140805
e-06
## At iteration 11 value of x is: 1.097153 0.9375543 and f is: -5.774754e-07 6.757782
e-07
## At iteration 12 value of x is: 1.097152 0.9375547 and f is: 1.822889e-07 -2.133196
e-07
## At iteration 13 value of x is: 1.097153 0.9375546 and f is: -5.754227e-08 6.733757
e-08
## At iteration 14 value of x is: 1.097153 0.9375546 and f is: 1.816409e-08 -2.125613
e-08
## At iteration 15 value of x is: 1.097153 0.9375546 and f is: -5.733772e-09 6.709819
e-09
## At iteration 16 value of x is: 1.097153 0.9375546 and f is: 1.809952e-09 -2.118057
e-09
## At iteration 17 value of x is: 1.097153 0.9375546 and f is: -5.713394e-10 6.685968
e-10
## Algorithm converged
```

```
llpois(res2[1:2], bs, N)
```

```
## [1] -48.02716
```

```
newtonmulti(gradpois, ipois, t(as.matrix(c(0.5,0.5))), bs, N)
```

```
## At iteration 1 value of x is: 1.117785 0.9062845 and f is: -0.05712725 0.07045913
## At iteration 2 value of x is: 1.090702 0.9473314 and f is: 0.01871967 -0.02155272
## At iteration 3 value of x is: 1.099195 0.9344591 and f is: -0.005841277 0.00687103
7
## At iteration 4 value of x is: 1.096508 0.9385308 and f is: 0.001850667 -0.00216218
## At iteration 5 value of x is: 1.097356 0.9372463 and f is: -0.0005835163 0.0006831
982
## At iteration 6 value of x is: 1.097088 0.9376519 and f is: 0.000184263 -0.00021559
48
## At iteration 7 value of x is: 1.097173 0.9375239 and f is: -5.815872e-05 6.806245e
-05
## At iteration 8 value of x is: 1.097146 0.9375643 and f is: 1.835935e-05 -2.148428e
-05
## At iteration 9 value of x is: 1.097155 0.9375515 and f is: -5.795342e-06 6.781905e
-06
## At iteration 10 value of x is: 1.097152 0.9375556 and f is: 1.829394e-06 -2.140805
e-06
## At iteration 11 value of x is: 1.097153 0.9375543 and f is: -5.774754e-07 6.757782
e-07
## At iteration 12 value of x is: 1.097152 0.9375547 and f is: 1.822889e-07 -2.133196
e-07
## At iteration 13 value of x is: 1.097153 0.9375546 and f is: -5.754227e-08 6.733757
e-08
## At iteration 14 value of x is: 1.097153 0.9375546 and f is: 1.816409e-08 -2.125613
e-08
## At iteration 15 value of x is: 1.097153 0.9375546 and f is: -5.733772e-09 6.709819
e-09
## At iteration 16 value of x is: 1.097153 0.9375546 and f is: 1.809952e-09 -2.118057
e-09
## At iteration 17 value of x is: 1.097153 0.9375546 and f is: -5.713394e-10 6.685968
e-10
## Algorithm converged
```

```
##          root1          root2          f.val1          f.val2
## 1.097153e+00  9.375546e-01 -5.713394e-10  6.685968e-10
```

```
newtonmulti(gradpois, ipois, t(as.matrix(c(0.1,0.5))), bs, N)
```



```

## At iteration 1 value of x is: 1.417701 0.451749 and f is: -0.4419727 1.38702
## At iteration 2 value of x is: 1.009561 1.070303 and f is: 0.2901992 -0.2737298
## At iteration 3 value of x is: 1.125967 0.8938844 and f is: -0.07863297 0.09904877
## At iteration 4 value of x is: 1.088177 0.9511568 and f is: 0.02615637 -0.02992437
## At iteration 5 value of x is: 1.099998 0.9332428 and f is: -0.008125141 0.00957696
6
## At iteration 6 value of x is: 1.096256 0.9389139 and f is: 0.002577974 -0.00300998
7
## At iteration 7 value of x is: 1.097436 0.9371253 and f is: -0.0008124677 0.0009514
534
## At iteration 8 value of x is: 1.097063 0.9376901 and f is: 0.0002565982 -0.0003002
105
## At iteration 9 value of x is: 1.097181 0.9375118 and f is: -8.098612e-05 9.477898e
-05
## At iteration 10 value of x is: 1.097144 0.9375681 and f is: 2.556579e-05 -2.991713
e-05
## At iteration 11 value of x is: 1.097155 0.9375503 and f is: -8.070102e-06 9.443926
e-06
## At iteration 12 value of x is: 1.097152 0.9375559 and f is: 2.547463e-06 -2.981106
e-06
## At iteration 13 value of x is: 1.097153 0.9375542 and f is: -8.041441e-07 9.410326
e-07
## At iteration 14 value of x is: 1.097152 0.9375547 and f is: 2.538404e-07 -2.970511
e-07
## At iteration 15 value of x is: 1.097153 0.9375545 and f is: -8.012857e-08 9.37687e
-08
## At iteration 16 value of x is: 1.097153 0.9375546 and f is: 2.52938e-08 -2.959951e
-08
## At iteration 17 value of x is: 1.097153 0.9375546 and f is: -7.984371e-09 9.343536
e-09
## At iteration 18 value of x is: 1.097153 0.9375546 and f is: 2.520388e-09 -2.949429
e-09
## At iteration 19 value of x is: 1.097153 0.9375546 and f is: -7.95598e-10 9.310322e
-10
## Algorithm converged

```

```

##          root1          root2          f.val1          f.val2
## 1.097153e+00  9.375546e-01 -7.955980e-10  9.310322e-10

```

```

newtonmulti(gradpois, ipois, t(as.matrix(c(0.8,0.5))), bs, N)

```

```

## At iteration 1 value of x is: 1.061965 0.9908827 and f is: 0.1071518 -0.1148385
## At iteration 2 value of x is: 1.108444 0.9204415 and f is: -0.03177974 0.03827084
## At iteration 3 value of x is: 1.093607 0.9429283 and f is: 0.01023817 -0.01187422
## At iteration 4 value of x is: 1.098274 0.9358555 and f is: -0.003211356 0.00376868
8
## At iteration 5 value of x is: 1.096799 0.9380907 and f is: 0.001015757 -0.00118760
5
## At iteration 6 value of x is: 1.097264 0.9373853 and f is: -0.0003204355 0.0003750
884
## At iteration 7 value of x is: 1.097117 0.937608 and f is: 0.0001011706 -0.00011838
21
## At iteration 8 value of x is: 1.097164 0.9375377 and f is: -3.193402e-05 3.737113e
-05
## At iteration 9 value of x is: 1.097149 0.9375599 and f is: 1.008066e-05 -1.179656e
-05
## At iteration 10 value of x is: 1.097154 0.9375529 and f is: -3.182093e-06 3.723784
e-06
## At iteration 11 value of x is: 1.097152 0.9375551 and f is: 1.004478e-06 -1.175467
e-06
## At iteration 12 value of x is: 1.097153 0.9375544 and f is: -3.170785e-07 3.710542
e-07
## At iteration 13 value of x is: 1.097152 0.9375546 and f is: 1.000907e-07 -1.171289
e-07
## At iteration 14 value of x is: 1.097153 0.9375546 and f is: -3.159514e-08 3.697351
e-08
## At iteration 15 value of x is: 1.097153 0.9375546 and f is: 9.973485e-09 -1.167125
e-08
## At iteration 16 value of x is: 1.097153 0.9375546 and f is: -3.148283e-09 3.684207
e-09
## At iteration 17 value of x is: 1.097153 0.9375546 and f is: 9.938049e-10 -1.162975
e-09
## At iteration 18 value of x is: 1.097153 0.9375546 and f is: -3.137112e-10 3.671102
e-10
## Algorithm converged

```

```

##          root1          root2          f.val1          f.val2
## 1.097153e+00  9.375546e-01 -3.137112e-10  3.671102e-10

```

```

newtonmulti(gradpois, ipois, t(as.matrix(c(2,0.5))), bs, N)

```

```

## At iteration 1 value of x is: 0.9969554 1.089408 and f is: 0.33856 -0.3098281
## At iteration 2 value of x is: 1.130314 0.8872973 and f is: -0.0897943 0.1143875
## At iteration 3 value of x is: 1.086844 0.9531772 and f is: 0.03010956 -0.03433193
## At iteration 4 value of x is: 1.100422 0.9325993 and f is: -0.009330961 0.01101008
## At iteration 5 value of x is: 1.096122 0.9391164 and f is: 0.002962824 -0.00345816
2
## At iteration 6 value of x is: 1.097478 0.9370613 and f is: -0.0009335322 0.0010933
45
## At iteration 7 value of x is: 1.09705 0.9377103 and f is: 0.0002948559 -0.00034495
9
## At iteration 8 value of x is: 1.097185 0.9375054 and f is: -9.305857e-05 0.0001089
087
## At iteration 9 value of x is: 1.097142 0.9375701 and f is: 2.937706e-05 -3.437697e
-05
## At iteration 10 value of x is: 1.097156 0.9375497 and f is: -9.273145e-06 1.085178
e-05
## At iteration 11 value of x is: 1.097152 0.9375561 and f is: 2.927226e-06 -3.425513
e-06
## At iteration 12 value of x is: 1.097153 0.9375541 and f is: -9.240217e-07 1.081317
e-06
## At iteration 13 value of x is: 1.097152 0.9375547 and f is: 2.916816e-07 -3.413339
e-07
## At iteration 14 value of x is: 1.097153 0.9375545 and f is: -9.207372e-08 1.077472
e-07
## At iteration 15 value of x is: 1.097153 0.9375546 and f is: 2.906447e-08 -3.401205
e-08
## At iteration 16 value of x is: 1.097153 0.9375546 and f is: -9.174642e-09 1.073642
e-08
## At iteration 17 value of x is: 1.097153 0.9375546 and f is: 2.896115e-09 -3.389114
e-09
## At iteration 18 value of x is: 1.097153 0.9375546 and f is: -9.142008e-10 1.069826
e-09
## At iteration 19 value of x is: 1.097153 0.9375546 and f is: 2.885815e-10 -3.377061
e-10
## Algorithm converged

```

```

##          root1          root2          f.val1          f.val2
## 1.097153e+00  9.375546e-01  2.885815e-10 -3.377061e-10

```

```

newtonmulti(gradpois, ipois, t(as.matrix(c(2,2))), bs, N)

```

```

## At iteration 1 value of x is: 1.117785 0.9062845 and f is: -0.05712725 0.07045913
## At iteration 2 value of x is: 1.090702 0.9473314 and f is: 0.01871967 -0.02155272
## At iteration 3 value of x is: 1.099195 0.9344591 and f is: -0.005841277 0.00687103
7
## At iteration 4 value of x is: 1.096508 0.9385308 and f is: 0.001850667 -0.00216218
## At iteration 5 value of x is: 1.097356 0.9372463 and f is: -0.0005835163 0.0006831
982
## At iteration 6 value of x is: 1.097088 0.9376519 and f is: 0.000184263 -0.00021559
48
## At iteration 7 value of x is: 1.097173 0.9375239 and f is: -5.815872e-05 6.806245e
-05
## At iteration 8 value of x is: 1.097146 0.9375643 and f is: 1.835935e-05 -2.148428e
-05
## At iteration 9 value of x is: 1.097155 0.9375515 and f is: -5.795342e-06 6.781905e
-06
## At iteration 10 value of x is: 1.097152 0.9375556 and f is: 1.829394e-06 -2.140805
e-06
## At iteration 11 value of x is: 1.097153 0.9375543 and f is: -5.774754e-07 6.757782
e-07
## At iteration 12 value of x is: 1.097152 0.9375547 and f is: 1.822889e-07 -2.133196
e-07
## At iteration 13 value of x is: 1.097153 0.9375546 and f is: -5.754227e-08 6.733757
e-08
## At iteration 14 value of x is: 1.097153 0.9375546 and f is: 1.816409e-08 -2.125613
e-08
## At iteration 15 value of x is: 1.097153 0.9375546 and f is: -5.733772e-09 6.709819
e-09
## At iteration 16 value of x is: 1.097153 0.9375546 and f is: 1.809952e-09 -2.118057
e-09
## At iteration 17 value of x is: 1.097153 0.9375546 and f is: -5.713394e-10 6.685968
e-10
## Algorithm converged

```

```

##          root1          root2          f.val1          f.val2
## 1.097153e+00  9.375546e-01 -5.713394e-10  6.685968e-10

```

All converge to the root found in (a),  $\hat{\alpha}_{MLE} = (1.0972, 0.9376)$ , even the (2,2) starting point which converged to a negative  $\alpha_2$  previously.

## (c)

Compare the implementation ease and performance of the two methods for this dataset.

The Fisher approach results in a matrix with slightly fewer terms in each position than the Hessian, but the derivation and implementation were fairly similar. We can look at runtime as a measure of performance, however:

```

tot_time_H <- numeric(100)
tot_time_F <- numeric(100)
for (i in 1:100) {
  # Run time to generate Hessian matrix
  start_time_H <- Sys.time()
  hesspois(testalphas, bs, N)
  end_time_H <- Sys.time()
  tot_time_H[i] <- as.numeric(end_time_H - start_time_H)
  #print(paste0("Hessian method takes ", tot_time_H))

  # Run time to generate Fisher Information matrix
  start_time_F <- Sys.time()
  ipois(testalphas, bs, N)
  end_time_F <- Sys.time()
  tot_time_F[i] <- as.numeric(end_time_F - start_time_F)
  #print(paste0("Fisher Information method takes ", tot_time_F))
}
print(paste0("Hessian method takes an average of ", mean(tot_time_H)," over 100 attempts on the same data."))

```

```

## [1] "Hessian method takes an average of 1.49726867675781e-05 over 100 attempts on the same data."

```

```

print(paste0("Fisher Information method takes an average of ", mean(tot_time_F)," over 100 attempts on the same data."))

```

```

## [1] "Fisher Information method takes an average of 1.29485130310059e-05 over 100 attempts on the same data."

```

```

print(paste0("Fisher Information method is, on average, ",round((mean(tot_time_H) - mean(tot_time_F))/mean(tot_time_H)*100, digits = 3),"% faster than the Hessian method on this data."))

```

```

## [1] "Fisher Information method is, on average, 13.519% faster than the Hessian method on this data."

```

We can see that calculating the Fisher Information matrix is slightly faster than calculating the Hessian.

## (d)

Estimate the standard errors of the MLEs of  $\alpha_1$  and  $\alpha_2$ .

We can estimate the standard errors by taking the square root of the diagonal of the inverse Hessian (or its approximation):

```
sqrt(diag(solve(-hesspois(res[1:2], bs, N)))) # Estimated standard errors of MLEs is square root of diagonal of inverse Hessian matrix
```

```
## [1] 0.3896129 0.5546760
```

```
sqrt(diag(solve(-ipois(res[1:2], bs, N)))) # Estimated standard errors of MLEs is square root of diagonal of inverse Fisher Information matrix
```

```
## [1] 0.4375560 0.6314687
```

We see that the standard errors when using the actual Hessian matrix are smaller than when using the Fisher Information matrix in its place. In both cases, the standard error for  $\alpha_1$  is smaller than that of  $\alpha_2$ .